BRACT's

Vishwakarma Institute of Information Technology, Pune Practical Implementation Sheet

Department: CS-AI Semester: 5 Academic Year: 2026-27

Division: E Batch: 2 Practical No:7 Roll no: 21 PRN No: 12411682

Subject Name: Deep Learning

Name of Student: Pranay Zunzunwala

Problem Statement: Implement transfer learning using pre-trained AlexNet, VGG16, ResNet50, and EfficientNetB0 models for image classification, and compare their performance.

Algorithm:

Step 1: Import the required libraries such as NumPy, Pandas, Matplotlib, TensorFlow/Keras, and the necessary modules for **AlexNet, VGG16, ResNet50, and EfficientNetB0**.

Step 2: Install and configure TensorFlow/Keras in Google Colab and verify the TensorFlow version and **GPU availability** for efficient model training.

Step 3: Load the **CIFAR-10 image classification dataset** using Keras and separate it into training and testing images along with their corresponding class labels.

Step 4: Inspect the dataset by checking the dimensions, number of classes, image shape, class labels, and displaying sample images from the dataset.

Step 5: Preprocess the image data by resizing the images to a common input size and applying the appropriate **normalization/preprocessing functions** required by the respective pretrained models.

Step 6: Create the **AlexNet architecture** and adapt its final classification layer for the 10 CIFAR-10 classes. For VGG16, ResNet50, and EfficientNetB0, load the **ImageNet-pretrained weights** and remove their original classification layers.

Step 7: Freeze the convolutional base of the pretrained **VGG16, ResNet50, and EfficientNetB0** models and add new classification layers consisting of Global Average Pooling, Dense, Dropout, and a Softmax output layer.

Step 8: Compile and train all four models using the **Adam optimizer**, sparse categorical cross-entropy loss, and accuracy as the evaluation metric, while recording their training time and validation performance.

Step 9: Evaluate each trained model using the test dataset and compare their **test accuracy, test loss, number of parameters, and training time**.

Step 10: Visualize and analyze the performance of AlexNet, VGG16, ResNet50, and EfficientNetB0 using **accuracy and loss comparison graphs, validation curves, and sample predictions**, and summarize the differences between the models.

Code:

```
# IMPORT LIBRARIES AND CHECK GPU

import os
import time
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf

from IPython.display import display
from tensorflow.keras import layers, models
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.applications import VGG16, ResNet50, EfficientNetB0
from tensorflow.keras.callbacks import EarlyStopping

print("TensorFlow version:", tf.__version__)
print("GPU devices:", tf.config.list_physical_devices("GPU"))

SEED = 42
random.seed(SEED)
np.random.seed(SEED)
tf.random.set_seed(SEED)

# Optional: avoid TensorFlow grabbing all GPU memory at once
gpus = tf.config.list_physical_devices("GPU")
for gpu in gpus:
    try:
        tf.config.experimental.set_memory_growth(gpu, True)
    except RuntimeError:
        pass

# CONFIGURATION

USE_FULL_DATASET = False   # True = all 50,000 train + 10,000 test images
IMG_SIZE = 128
BATCH_SIZE = 32
EPOCHS = 5
NUM_CLASSES = 10

if USE_FULL_DATASET:
```

```
    TRAIN_LIMIT = None
    TEST_LIMIT = None
else:
    TRAIN_LIMIT = 10000
    TEST_LIMIT = 2000

print("Configuration:")
print("Image size:", IMG_SIZE)
print("Batch size:", BATCH_SIZE)
print("Epochs:", EPOCHS)
print("Full dataset:", USE_FULL_DATASET)

# LOAD CIFAR-10 DATASET

(x_train, y_train), (x_test, y_test) = cifar10.load_data()

if TRAIN_LIMIT is not None:
    x_train = x_train[:TRAIN_LIMIT]
    y_train = y_train[:TRAIN_LIMIT]

if TEST_LIMIT is not None:
    x_test = x_test[:TEST_LIMIT]
    y_test = y_test[:TEST_LIMIT]

y_train = y_train.flatten()
y_test = y_test.flatten()

class_names = [
    "Airplane", "Automobile", "Bird", "Cat", "Deer",
    "Dog", "Frog", "Horse", "Ship", "Truck"
]

print("Training images:", x_train.shape)
print("Training labels:", y_train.shape)
print("Testing images:", x_test.shape)
print("Testing labels:", y_test.shape)

# DISPLAY SAMPLE IMAGES

plt.figure(figsize=(12, 8))

for i in range(20):
    plt.subplot(4, 5, i + 1)
    plt.imshow(x_train[i])
    plt.title(class_names[y_train[i]])
    plt.axis("off")
```

```
plt.tight_layout()
plt.show()

# CREATE TF.DATA INPUT PIPELINE

def resize_images(image, label):
    image = tf.image.resize(image, (IMG_SIZE, IMG_SIZE))
    image = tf.cast(image, tf.float32)
    return image, label

train_ds = tf.data.Dataset.from_tensor_slices((x_train, y_train))
test_ds = tf.data.Dataset.from_tensor_slices((x_test, y_test))

train_ds = (
    train_ds
    .shuffle(min(len(x_train), 10000), seed=SEED)
    .map(resize_images, num_parallel_calls=tf.data.AUTOTUNE)
    .batch(BATCH_SIZE)
    .prefetch(tf.data.AUTOTUNE)
)

test_ds = (
    test_ds
    .map(resize_images, num_parallel_calls=tf.data.AUTOTUNE)
    .batch(BATCH_SIZE)
    .prefetch(tf.data.AUTOTUNE)
)

print("Dataset pipeline created successfully.")

## AlexNet

# CREATE ALEXNET

def create_alexnet():
    model = models.Sequential([
        layers.Input(shape=(IMG_SIZE, IMG_SIZE, 3)),
        layers.Rescaling(1.0 / 255.0),

        layers.Conv2D(96, kernel_size=11, strides=4, activation="relu"),
        layers.BatchNormalization(),
        layers.MaxPooling2D(pool_size=3, strides=2),

        layers.Conv2D(256, kernel_size=5, padding="same", activation="relu"),
        layers.BatchNormalization(),
```

```
        layers.MaxPooling2D(pool_size=3, strides=2),

        layers.Conv2D(384, kernel_size=3, padding="same", activation="relu"),
        layers.Conv2D(384, kernel_size=3, padding="same", activation="relu"),
        layers.Conv2D(256, kernel_size=3, padding="same", activation="relu"),

        layers.MaxPooling2D(pool_size=3, strides=2),

        layers.Flatten(),
        layers.Dense(512, activation="relu"),
        layers.Dropout(0.5),
        layers.Dense(256, activation="relu"),
        layers.Dropout(0.5),

        layers.Dense(NUM_CLASSES, activation="softmax")
    ])

    model.compile(
        optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
        loss="sparse_categorical_crossentropy",
        metrics=["accuracy"]
    )

    return model

alexnet = create_alexnet()
alexnet.summary()

# TRAIN ALEXNET

early_stopping_alexnet = EarlyStopping(
    monitor="val_loss",
    patience=2,
    restore_best_weights=True
)

start = time.time()

history_alexnet = alexnet.fit(
    train_ds,
    validation_data=test_ds,
    epochs=EPOCHS,
    callbacks=[early_stopping_alexnet],
    verbose=1
)
```

```
alexnet_time = time.time() - start

alexnet_loss, alexnet_accuracy = alexnet.evaluate(test_ds, verbose=0)

print(f"AlexNet test accuracy: {alexnet_accuracy:.4f}")
print(f"AlexNet test loss: {alexnet_loss:.4f}")
print(f"AlexNet training time: {alexnet_time/60:.2f} minutes")

# CREATE TRANSFER LEARNING MODEL

def create_transfer_model(model_name):
    if model_name == "VGG16":
        base_model = VGG16(
            weights="imagenet",
            include_top=False,
            input_shape=(IMG_SIZE, IMG_SIZE, 3)
        )
        preprocess_input = tf.keras.applications.vgg16.preprocess_input

    elif model_name == "ResNet50":
        base_model = ResNet50(
            weights="imagenet",
            include_top=False,
            input_shape=(IMG_SIZE, IMG_SIZE, 3)
        )
        preprocess_input = tf.keras.applications.resnet50.preprocess_input

    elif model_name == "EfficientNetB0":
        base_model = EfficientNetB0(
            weights="imagenet",
            include_top=False,
            input_shape=(IMG_SIZE, IMG_SIZE, 3)
        )
        preprocess_input = tf.keras.applications.efficientnet.preprocess_input

    else:
        raise ValueError("Supported models: VGG16, ResNet50, EfficientNetB0")

    # Freeze pretrained feature extractor
    base_model.trainable = False

    inputs = layers.Input(shape=(IMG_SIZE, IMG_SIZE, 3))

    x = preprocess_input(inputs)
    x = base_model(x, training=False)
    x = layers.GlobalAveragePooling2D()(x)
```

```
    x = layers.Dense(256, activation="relu")(x)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(NUM_CLASSES, activation="softmax")(x)

    model = models.Model(inputs, outputs, name=model_name)

    model.compile(
        optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
        loss="sparse_categorical_crossentropy",
        metrics=["accuracy"]
    )

    return model

# TRAIN VGG16

vgg16 = create_transfer_model("VGG16")
vgg16.summary()

early_stopping_vgg = EarlyStopping(
    monitor="val_loss",
    patience=2,
    restore_best_weights=True
)

start = time.time()

history_vgg16 = vgg16.fit(
    train_ds,
    validation_data=test_ds,
    epochs=EPOCHS,
    callbacks=[early_stopping_vgg],
    verbose=1
)

vgg_time = time.time() - start

vgg_loss, vgg_accuracy = vgg16.evaluate(test_ds, verbose=0)

print(f"VGG16 test accuracy: {vgg_accuracy:.4f}")
print(f"VGG16 test loss: {vgg_loss:.4f}")
print(f"VGG16 training time: {vgg_time/60:.2f} minutes")

# TRAIN RESNET50

resnet50 = create_transfer_model("ResNet50")
```

```
resnet50.summary()

early_stopping_resnet = EarlyStopping(
    monitor="val_loss",
    patience=2,
    restore_best_weights=True
)

start = time.time()

history_resnet = resnet50.fit(
    train_ds,
    validation_data=test_ds,
    epochs=EPOCHS,
    callbacks=[early_stopping_resnet],
    verbose=1
)

resnet_time = time.time() - start

resnet_loss, resnet_accuracy = resnet50.evaluate(test_ds, verbose=0)

print(f"ResNet50 test accuracy: {resnet_accuracy:.4f}")
print(f"ResNet50 test loss: {resnet_loss:.4f}")
print(f"ResNet50 training time: {resnet_time/60:.2f} minutes")

# TRAIN EFFICIENTNETB0

efficientnet = create_transfer_model("EfficientNetB0")
efficientnet.summary()

early_stopping_eff = EarlyStopping(
    monitor="val_loss",
    patience=2,
    restore_best_weights=True
)

start = time.time()

history_efficientnet = efficientnet.fit(
    train_ds,
    validation_data=test_ds,
    epochs=EPOCHS,
    callbacks=[early_stopping_eff],
    verbose=1
)
```

```
efficientnet_time = time.time() - start

efficientnet_loss, efficientnet_accuracy = efficientnet.evaluate(test_ds, verbose=0)

print(f"EfficientNetB0 test accuracy: {efficientnet_accuracy:.4f}")
print(f"EfficientNetB0 test loss: {efficientnet_loss:.4f}")
print(f"EfficientNetB0 training time: {efficientnet_time/60:.2f} minutes")

# PERFORMANCE COMPARISON TABLE

results = pd.DataFrame({
    "Model": [
        "AlexNet",
        "VGG16",
        "ResNet50",
        "EfficientNetB0"
    ],
    "Test Accuracy": [
        alexnet_accuracy,
        vgg_accuracy,
        resnet_accuracy,
        efficientnet_accuracy
    ],
    "Test Loss": [
        alexnet_loss,
        vgg_loss,
        resnet_loss,
        efficientnet_loss
    ],
    "Training Time (min)": [
        alexnet_time / 60,
        vgg_time / 60,
        resnet_time / 60,
        efficientnet_time / 60
    ],
    "Parameters": [
        alexnet.count_params(),
        vgg16.count_params(),
        resnet50.count_params(),
        efficientnet.count_params()
    ]
})

results = results.sort_values(
    by="Test Accuracy",
```

```
    ascending=False
).reset_index(drop=True)

display(results)

# Save results to CSV
results.to_csv("transfer_learning_results.csv", index=False)
print("Results saved as transfer_learning_results.csv")

#  ACCURACY COMPARISON

plt.figure(figsize=(10, 6))

bars = plt.bar(
    results["Model"],
    results["Test Accuracy"]
)

plt.xlabel("Model")
plt.ylabel("Test Accuracy")
plt.title("Test Accuracy Comparison")
plt.ylim(0, 1)

for bar, value in zip(bars, results["Test Accuracy"]):
    plt.text(
        bar.get_x() + bar.get_width() / 2,
        value + 0.01,
        f"{value:.4f}",
        ha="center"
    )

plt.tight_layout()
plt.show()

# TEST LOSS COMPARISON

plt.figure(figsize=(10, 6))

bars = plt.bar(
    results["Model"],
    results["Test Loss"]
)

plt.xlabel("Model")
plt.ylabel("Test Loss")
plt.title("Test Loss Comparison")
```

```
for bar, value in zip(bars, results["Test Loss"]):
    plt.text(
        bar.get_x() + bar.get_width() / 2,
        value,
        f"{value:.3f}",
        ha="center",
        va="bottom"
    )

plt.tight_layout()
plt.show()

#  TRAINING TIME COMPARISON

plt.figure(figsize=(10, 6))

bars = plt.bar(
    results["Model"],
    results["Training Time (min)"]
)

plt.xlabel("Model")
plt.ylabel("Training Time (minutes)")
plt.title("Training Time Comparison")

for bar, value in zip(bars, results["Training Time (min)"]):
    plt.text(
        bar.get_x() + bar.get_width() / 2,
        value,
        f"{value:.2f}",
        ha="center",
        va="bottom"
    )

plt.tight_layout()
plt.show()

#  VALIDATION ACCURACY CURVES

plt.figure(figsize=(11, 7))

plt.plot(history_alexnet.history["val_accuracy"], label="AlexNet")
plt.plot(history_vgg16.history["val_accuracy"], label="VGG16")
plt.plot(history_resnet.history["val_accuracy"], label="ResNet50")
plt.plot(history_efficientnet.history["val_accuracy"], label="EfficientNetB0")
```

```
plt.xlabel("Epoch")
plt.ylabel("Validation Accuracy")
plt.title("Validation Accuracy Comparison")
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()

# VALIDATION LOSS CURVES

plt.figure(figsize=(11, 7))

plt.plot(history_alexnet.history["val_loss"], label="AlexNet")
plt.plot(history_vgg16.history["val_loss"], label="VGG16")
plt.plot(history_resnet.history["val_loss"], label="ResNet50")
plt.plot(history_efficientnet.history["val_loss"], label="EfficientNetB0")

plt.xlabel("Epoch")
plt.ylabel("Validation Loss")
plt.title("Validation Loss Comparison")
plt.legend()
plt.grid(True)

plt.tight_layout()
plt.show()

# SINGLE IMAGE PREDICTION

index = 10

original_image = x_test[index]
actual_label = y_test[index]

image_resized = tf.image.resize(
    original_image,
    (IMG_SIZE, IMG_SIZE)
)

image_batch = tf.expand_dims(image_resized, axis=0)

prediction = efficientnet.predict(image_batch, verbose=0)

predicted_label = int(np.argmax(prediction[0]))
confidence = float(np.max(prediction[0]))
```

```
plt.figure(figsize=(5, 5))
plt.imshow(original_image)
plt.title(
    f"Actual: {class_names[actual_label]}\n"
    f"Predicted: {class_names[predicted_label]} "
    f"({confidence * 100:.2f}%)"
)
plt.axis("off")
plt.show()

print("Actual class:", class_names[actual_label])
print("Predicted class:", class_names[predicted_label])
print(f"Confidence: {confidence * 100:.2f}%")
```

Output:

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 128, 128, 3)	0
conv2d (Conv2D)	(None, 30, 30, 96)	34,944
batch_normalization (BatchNormalization)	(None, 30, 30, 96)	384
max_pooling2d (MaxPooling2D)	(None, 14, 14, 96)	0
conv2d_1 (Conv2D)	(None, 14, 14, 256)	614,656
batch_normalization_1 (BatchNormalization)	(None, 14, 14, 256)	1,024
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 256)	0
conv2d_2 (Conv2D)	(None, 6, 6, 384)	885,120
conv2d_3 (Conv2D)	(None, 6, 6, 384)	1,327,488
conv2d_4 (Conv2D)	(None, 6, 6, 256)	884,992
max_pooling2d_2 (MaxPooling2D)	(None, 2, 2, 256)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 512)	524,800
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131,328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 10)	2,570

```
Total params: 4,407,306 (16.81 MB)

Trainable params: 4,406,602 (16.81 MB)

Non-trainable params: 704 (2.75 KB)
```

```
Epoch 1/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 432s 1s/step - accuracy: 0.2179 - loss: 2.1025 - val_accuracy: 0.1275 - val_loss: 2.8318
Epoch 2/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 445s 1s/step - accuracy: 0.2916 - loss: 1.8722 - val_accuracy: 0.3000 - val_loss: 1.8773
Epoch 3/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 459s 1s/step - accuracy: 0.3518 - loss: 1.7460 - val_accuracy: 0.3230 - val_loss: 1.8649
Epoch 4/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 488s 2s/step - accuracy: 0.3975 - loss: 1.6423 - val_accuracy: 0.4170 - val_loss: 1.6490
Epoch 5/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 432s 1s/step - accuracy: 0.4346 - loss: 1.5682 - val_accuracy: 0.3625 - val_loss: 1.8100
AlexNet test accuracy: 0.4170
AlexNet test loss: 1.6490
AlexNet training time: 37.78 minutes
```

Model: "ResNet50"

Layer (type)	Output Shape	Param #	Connected to
input_layer_4 (InputLayer)	(None, 128, 128, 3)	0	-
get_item_3 (GetItem)	(None, 128, 128)	0	input_layer_4[0]…
get_item_4 (GetItem)	(None, 128, 128)	0	input_layer_4[0]…
get_item_5 (GetItem)	(None, 128, 128)	0	input_layer_4[0]…
stack_1 (Stack)	(None, 128, 128, 3)	0	get_item_3[0][0], get_item_4[0][0], get_item_5[0][0]
add_1 (Add)	(None, 128, 128, 3)	0	stack_1[0][0]
resnet50 (Functional)	(None, 4, 4, 2048)	23,587,712	add_1[0][0]
global_average_poo… (GlobalAveragePool…	(None, 2048)	0	resnet50[0][0]
dense_5 (Dense)	(None, 256)	524,544	global_average_p…
dropout_3 (Dropout)	(None, 256)	0	dense_5[0][0]
dense_6 (Dense)	(None, 10)	2,570	dropout_3[0][0]

Total params: 24,114,826 (91.99 MB)

Trainable params: 527,114 (2.01 MB)

Non-trainable params: 23,587,712 (89.98 MB)

```
Epoch 1/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 738s 2s/step - accuracy: 0.7573 - loss: 0.7617 - val_accuracy: 0.8400 - val_loss: 0.4641
Epoch 2/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 739s 2s/step - accuracy: 0.8342 - loss: 0.4842 - val_accuracy: 0.8650 - val_loss: 0.3970
Epoch 3/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 723s 2s/step - accuracy: 0.8607 - loss: 0.4052 - val_accuracy: 0.8670 - val_loss: 0.3706
Epoch 4/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 698s 2s/step - accuracy: 0.8836 - loss: 0.3383 - val_accuracy: 0.8660 - val_loss: 0.3989
Epoch 5/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 768s 2s/step - accuracy: 0.8958 - loss: 0.3102 - val_accuracy: 0.8705 - val_loss: 0.4006
ResNet50 test accuracy: 0.8670
ResNet50 test loss: 0.3706
ResNet50 training time: 61.10 minutes
```

```
Model: "EfficientNetB0"
```

Layer (type)	Output Shape	Param #
input_layer_6 (InputLayer)	(None, 128, 128, 3)	0
efficientnetb0 (Functional)	(None, 4, 4, 1280)	4,049,571
global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 1280)	0
dense_7 (Dense)	(None, 256)	327,936
dropout_4 (Dropout)	(None, 256)	0
dense_8 (Dense)	(None, 10)	2,570

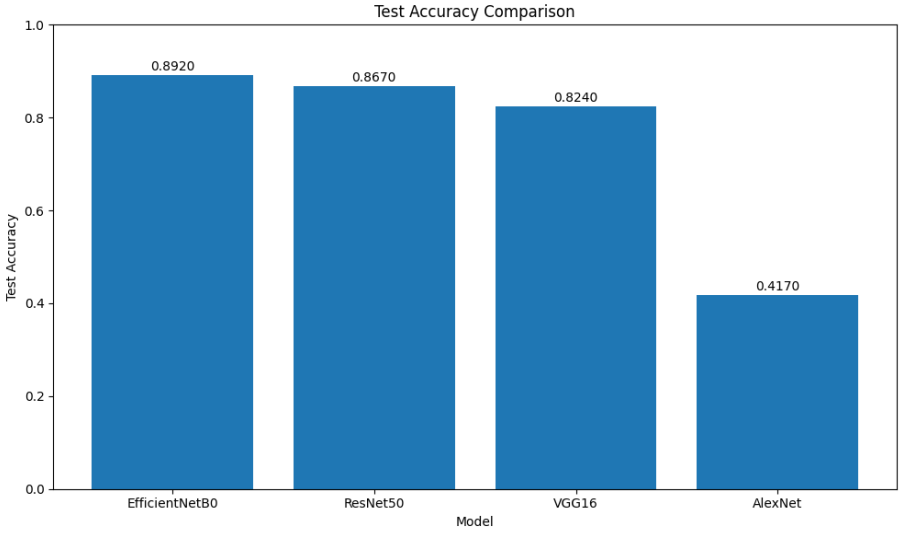
```
Total params: 4,380,077 (16.71 MB)

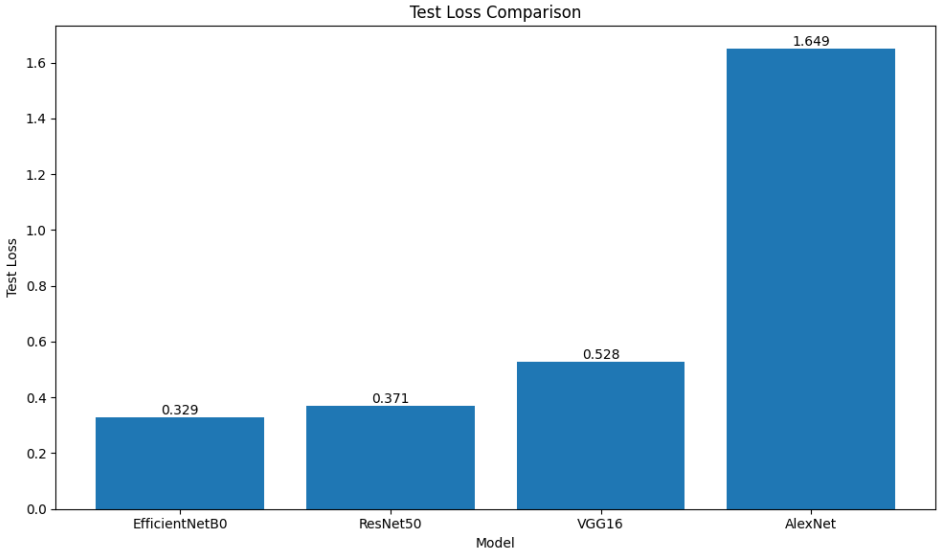
Trainable params: 330,506 (1.26 MB)

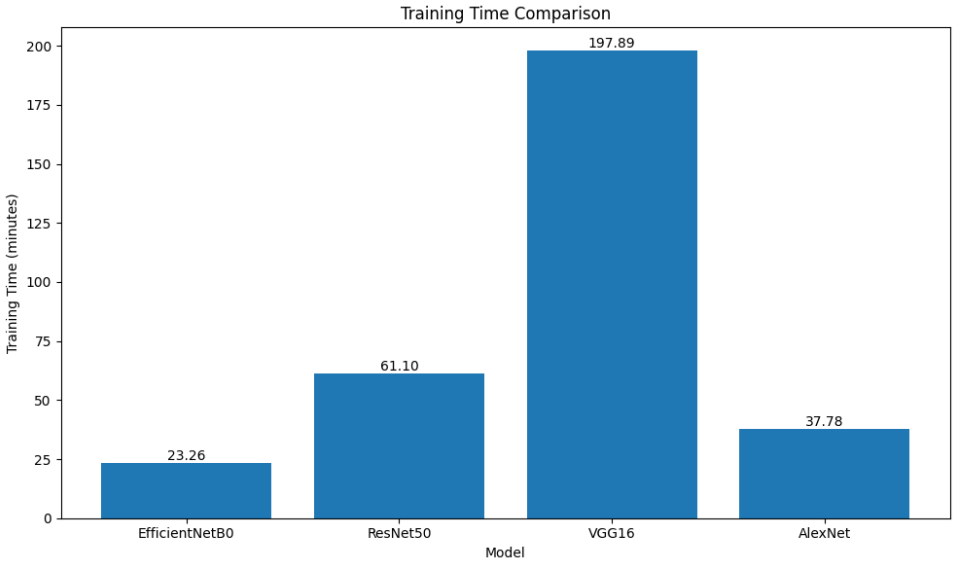
Non-trainable params: 4,049,571 (15.45 MB)

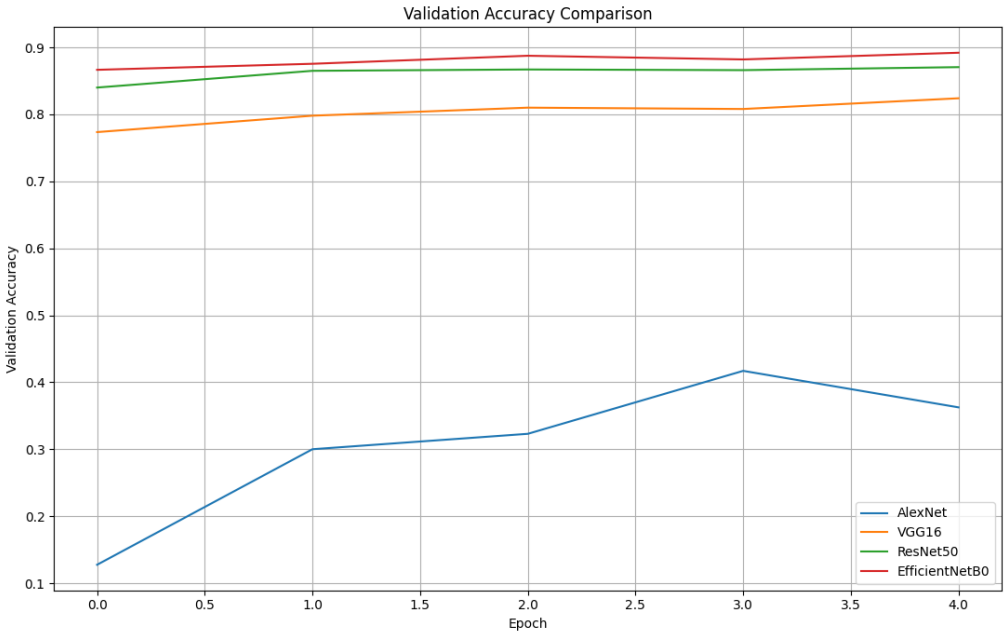
Epoch 1/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 274s 841ms/step - accuracy: 0.7977 - loss: 0.6241 - val_accuracy: 0.8665 - val_loss: 0.3744
Epoch 2/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 261s 833ms/step - accuracy: 0.8710 - loss: 0.3880 - val_accuracy: 0.8755 - val_loss: 0.3589
Epoch 3/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 261s 834ms/step - accuracy: 0.8919 - loss: 0.3217 - val_accuracy: 0.8875 - val_loss: 0.3341
Epoch 4/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 299s 957ms/step - accuracy: 0.9095 - loss: 0.2706 - val_accuracy: 0.8820 - val_loss: 0.3452
Epoch 5/5
313/313 ━━━━━━━━━━━━━━━━━━━━ 300s 961ms/step - accuracy: 0.9169 - loss: 0.2419 - val_accuracy: 0.8920 - val_loss: 0.3287
EfficientNetB0 test accuracy: 0.8920
EfficientNetB0 test loss: 0.3287
EfficientNetB0 training time: 23.26 minutes
```

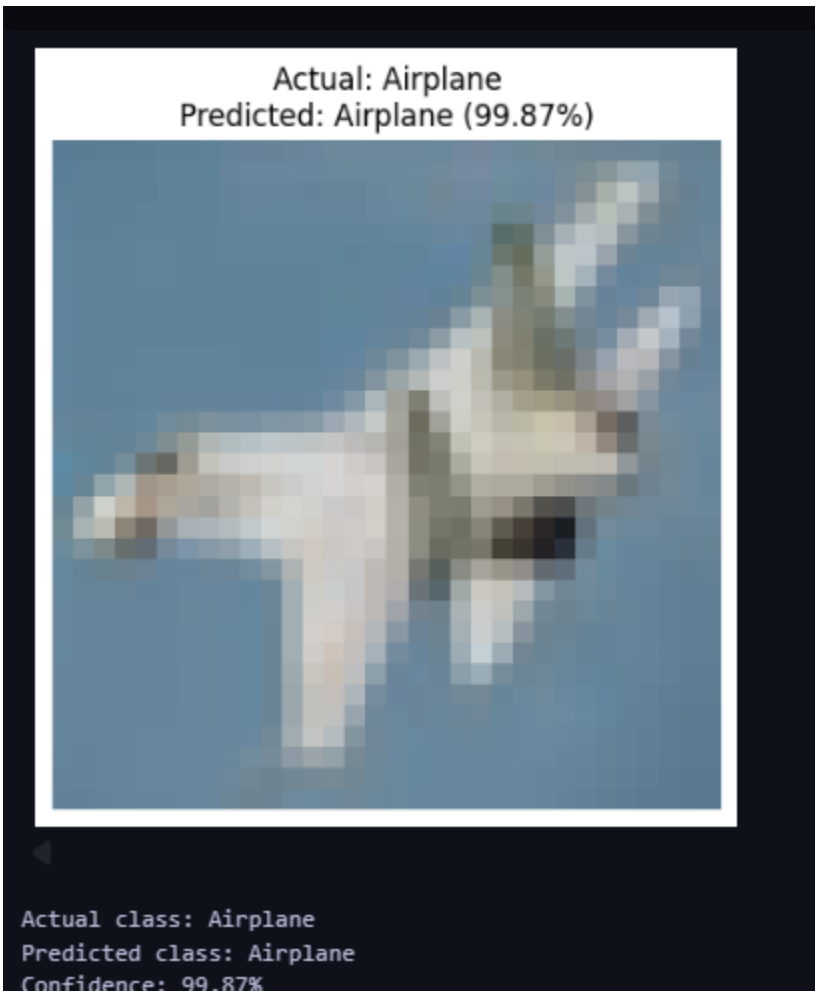
	Model	Test Accuracy	Test Loss	Training Time (min)	Parameters
0	EfficientNetB0	0.892	0.328724	23.260103	4380077
1	ResNet50	0.867	0.370568	61.097314	24114826
2	VGG16	0.824	0.528155	197.894695	14848586
3	AlexNet	0.417	1.649013	37.775874	4407306


Test Accuracy Comparison
Test Accuracy
Model
EfficientNetB0
ResNet50
VGG16
AlexNet
0.8920
0.8670
0.8240
0.4170
0.0
0.2
0.4
0.6
0.8
1.0


Test Loss Comparison
Test Loss
Model
EfficientNetB0
ResNet50
VGG16
AlexNet
0.329
0.371
0.528
1.649
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.4
1.6


Training Time Comparison
Training Time (minutes)
Model
EfficientNetB0
ResNet50
VGG16
AlexNet
23.26
61.10
197.89
37.78
0
25
50
75
100
125
150
175
200


Validation Accuracy Comparison
Validation Accuracy
Epoch
AlexNet
VGG16
ResNet50
EfficientNetB0


Actual: Airplane
Predicted: Airplane (99.87%)
Actual class: Airplane
Predicted class: Airplane
Confidence: 99.87%